

**National University of Singapore**  
**College of Design & Engineering – ECE**

**EE4400 Data Engineering and Deep Learning**  
**Tutorial 4 – CNN Solutions**

**Q1.** The convolution operation performs feature extraction through a filtering operation on the input values. The kernel or weights of the filter are learnt through back propagation, and different filters are created to extract different features from the input image, at low, medium and high levels in subsequent convolution layers.

**Q2.** Each filter acts as a receptive field and has its own set of weights/parameters. Each filter slides across the input image looking for certain features. This enables the filter to be invariant to the location of the feature, i.e. the feature can be at any location in the image. The weights/parameters related to a filter and receptive field can be viewed as being shared across the entire image.

**Q3.** (a) Use the  $(N-F)/\text{stride}+1$  formula with  $N=9$ ,  $F=5$  and  $\text{stride}=2$  to get  $3 \times 3$  pixels.

(b) The input image can be padded with 1 pixel around its 4 borders. Use the  $(N-F+2P)/\text{stride}+1$  formula to get  $4 \times 4$  pixels.

(c) No. of weights:  $F \times F + 1 = 26$ . Same in both cases.

**Q4.** Refer to the Colab notebook with results and explanations.

To understand how to obtain the numbers of parameters at each layer, refer to the lecture slides.

## ▼ EE4400 - Tutorial 4 Q4 Solution

This tutorial is on handwritten digit image classification using a convolutional neural network (CNN).

```
# To determine which version you're using:  
!pip show tensorflow  
  
Name: tensorflow  
Version: 2.4.1  
Summary: TensorFlow is an open source machine learning framework for everyone.  
Home-page: https://www.tensorflow.org/  
Author: Google Inc.  
Author-email: packages@tensorflow.org  
License: Apache 2.0  
Location: /usr/local/lib/python3.7/dist-packages  
Requires: tensorboard, absl-py, numpy, gast, google-pasta, h5py, protobuf, typing-extensions, wrapt, wheel,  
Required-by: fancyimpute
```

We shall use the Tensorflow package which contains the Keras sub-package to implement the CNN. The dataset is the MNIST dataset found within Keras. Import the necessary packages.

Refer to <https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-from-scratch-for-mnist-handwritten-digit-classification/>

```
## baseline cnn model for mnist  
from numpy import mean  
from numpy import std  
from matplotlib import pyplot  
from sklearn.model_selection import KFold  
# from tensorflow import keras  
# from tensorflow.keras import layers  
from tensorflow.keras.datasets import mnist  
from tensorflow.keras.utils import to_categorical  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Conv2D  
from tensorflow.keras.layers import MaxPooling2D  
from tensorflow.keras.layers import Dense  
from tensorflow.keras.layers import Flatten  
from tensorflow.keras.optimizers import SGD  
import warnings
```

Write a function to load the training and test datasets. Reshape the inputs to be in a single channel, and convert the outputs to one-hot-encoded format.

```
## load train and test dataset  
def load_dataset():  
    # load dataset  
    (trainX, trainY), (testX, testY) = mnist.load_data()  
    # reshape dataset to have a single channel  
    trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))  
    testX = testX.reshape((testX.shape[0], 28, 28, 1))  
    # one hot encode target values  
    trainY = to_categorical(trainY)  
    testY = to_categorical(testY)  
    return trainX, trainY, testX, testY
```

Write a function to prepare (convert integers to floats) and scale pixels (so that range is [0,1]) in the input images.

```
## scale pixels
def prep_pixels(train, test):
    # convert from integers to floats
    train_norm = train.astype('float32')
    test_norm = test.astype('float32')
    # normalize to range 0-1
    train_norm = train_norm / 255.0
    test_norm = test_norm / 255.0
    # return normalized images
    return train_norm, test_norm
```

Write a function `define_model()` to define the CNN model. Set up a CNN network with 16 filters each of size (3,3) followed by a 2D max pooling layer of pooling size (2,2). Add a dense layer of 50 nodes before the output with 10 nodes (i.e., 1 node per digit). Use the SGD optimizer and the categorical cross entropy loss function.

```
## define cnn model
def define_model():
    model = Sequential()
    model.add(Conv2D(16, (3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(28, 28, 1)))
    model.add(MaxPooling2D((2, 2)))
    model.add(Flatten())
    model.add(Dense(50, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dense(10, activation='softmax'))
    # compile model
    opt = SGD(lr=0.01, momentum=0.9)
    model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
    return model
```

Write a function `'evaluate_model(dataX, dataY, n_folds=5)`' to evaluate a model through n-fold cross-validation using the `sklearn.model_selection` function `'KFold()'`. The `define_model()`, `model.fit()` and `model.evaluate()` functions are used. The `'evaluate_model()'` function returns the 'scores' and 'histories', which are lists of accuracy values and model statistics, respectively.

```
## evaluate a model using k-fold cross-validation
def evaluate_model(dataX, dataY, n_folds=5):
    scores, histories = list(), list()
    # prepare cross validation
    kfold = KFold(n_folds, shuffle=True, random_state=1)
    # enumerate splits
    for train_ix, test_ix in kfold.split(dataX):
        # define model
        model = define_model()
        model.summary()
        # select rows for train and test
        trainX, trainY, testX, testY = dataX[train_ix], dataY[train_ix], dataX[test_ix], dataY[test_ix]
        # fit model
        history = model.fit(trainX, trainY, epochs=10, batch_size=32, validation_data=(testX, testY), verbose=0)
        # evaluate model
        _, acc = model.evaluate(testX, testY, verbose=0)
        print('> %.3f' % (acc * 100.0))
        # stores scores
        scores.append(acc)
        histories.append(history)
    return scores, histories
```

Write a function to plot the learning curves of cross entropy loss and classification accuracy, using 'histories' as input.

```
## plot diagnostic learning curves
def summarize_diagnostics(histories):
    for i in range(len(histories)):
        # plot loss
        pyplot.subplot(2, 1, 1)
```

```

pyp      p  ( , , )
pyplot.title('Cross Entropy Loss')
pyplot.plot(histories[i].history['loss'], color='blue', label='train')
pyplot.plot(histories[i].history['val_loss'], color='orange', label='test')
# plot accuracy
pyplot.subplot(2, 1, 2)
pyplot.title('Classification Accuracy')
pyplot.plot(histories[i].history['accuracy'], color='blue', label='train')
pyplot.plot(histories[i].history['val_accuracy'], color='orange', label='test')
pyplot.show()

```

Write a function to summarize the model performance using a box plot of the accuracy values, using 'scores' as input.

```

## summarize model performance
def summarize_performance(scores):
    # print summary
    print('Accuracy: mean=% .3f std=% .3f, n=%d' % (mean(scores)*100, std(scores)*100, len(scores)))
    # box and whisker plots of results
    pyplot.boxplot(scores)
    pyplot.show()

```

Write the main processing steps that calls the functions defined above to do the following steps: (i) load the dataset, (ii) prepare pixel data, (iii) evaluate model, (iv) plot learning curves, and (v) summarize estimated performance. Note: the evaluate\_model() function may take 5 to 10 minutes to run since the network and data are fairly large, and a lot of processing is required.

```

# suppress warnings
warnings.filterwarnings("ignore", category=UserWarning)

## main processing

# load dataset
trainX, trainY, testX, testY = load_dataset()

# prepare pixel data
trainX, testX = prep_pixels(trainX, testX)

print(trainX.shape)
print(trainY.shape)
print(testX.shape)
print(testY.shape)

```

```

(60000, 28, 28, 1)
(60000, 10)
(10000, 28, 28, 1)
(10000, 10)

```

```

# evaluate model
scores, histories = evaluate_model(trainX, trainY)
print(scores)
print(histories)

```

```
Model: "sequential_17"
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_27 (Conv2D)	(None, 26, 26, 16)	160
<hr/>		
max_pooling2d_27 (MaxPooling)	(None, 13, 13, 16)	0
<hr/>		
flatten_17 (Flatten)	(None, 2704)	0
<hr/>		
dense_34 (Dense)	(None, 50)	135250
<hr/>		
dense_35 (Dense)	(None, 10)	510

```
=====
Total params: 135,920
Trainable params: 135,920
Non-trainable params: 0
```

---

> 98.383  
Model: "sequential\_18"

Layer (type)	Output Shape	Param #
conv2d_28 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_28 (MaxPooling)	(None, 13, 13, 16)	0
flatten_18 (Flatten)	(None, 2704)	0
dense_36 (Dense)	(None, 50)	135250
dense_37 (Dense)	(None, 10)	510

---

```
Total params: 135,920
Trainable params: 135,920
Non-trainable params: 0
```

---

> 98.117  
Model: "sequential\_19"

Layer (type)	Output Shape	Param #
conv2d_29 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_29 (MaxPooling)	(None, 13, 13, 16)	0
flatten_19 (Flatten)	(None, 2704)	0
dense_38 (Dense)	(None, 50)	135250
dense_39 (Dense)	(None, 10)	510

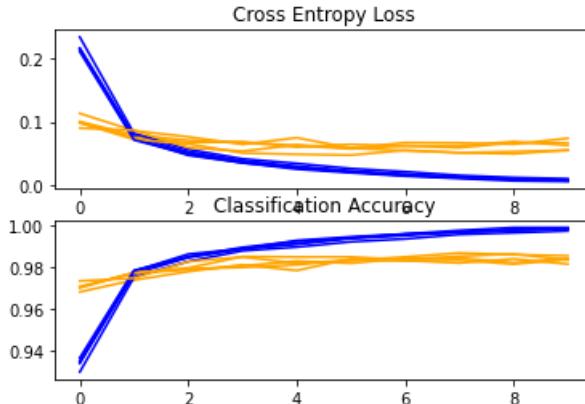
---

```
Total params: 135,920
Trainable params: 135,920
Non-trainable params: 0
```

---

> 98.333  
Model: "sequential\_20"

```
# learning curves
summarize_diagnostics(histories)
```



```
# summarize estimated performance
summarize_performance(scores)
```

Accuracy: mean=98.333 std=0.129, n=5

